Lecture 6: Designing a URL **Shortener like Bit.ly**

Functional & Non-Functional Requirements

- Functional Requirements
 - Accept long URL from user
 - Example: www.example.com/api/v3/userid=1234user_id=harshal
 - Generate short URL
 - Example: bit.ly/3xYzAb
 - Redirect short URL to original long URL

Non-Functional Requirements

- Scalability (handle millions of URLs)
- High availability & reliability
- Low latency (fast redirects)

High-Level Design (HLD)

Open-ended question → Discussion on requirements, functionality, and applications

URL Redirection Flow :

 $\texttt{Long URL} \ \rightarrow \ \texttt{Bit.ly (shortener)} \ \rightarrow \ \texttt{Short URL}$

Short URL \rightarrow Browser \rightarrow Redirect to Long URL

? Key Questions & Answers

Question	Answer
How many URLs to target?	20 million URLs
Example of a long URL	www.example.com/api/v3/userid=1234user_id=harshal
Now short should the short URL be?	As short as possible
Allowed characters in short URL	0-9, a-z, A-Z

(S) User Interaction Flow

- 1. **Q** User submits **long URL**
- 3. O User enters **short URL** in browser
- 4. System redirects to long URL

% Core Tasks

- Redirect short URL to long URL

Back of the Envelope Calculations – URL Shortener Design



4 1. Write QPS (Queries Per Second)



20 million URLs are generated per day

Calculation:

```
Write QPS = Total writes per day / Seconds per day
            = 20,000,000 / 86,400
            = (20 \times 10^{6}) / (86.4 \times 10^{3})
            = 0.0002314815 \times 10^{3}
            ≈ 231.48
```

Result:

Write QPS ≈ 232 writes/sec



2. Read QPS Calculation



Read to Write ratio = 10:1 (i.e., every short URL is read ~10 times)

Calculation:

```
Read QPS = Write QPS \times 10 = 232 \times 10 = 2,320 reads/sec
```

Note: Every read results in a database lookup

Result:

Read QPS ≈ 2,320 reads/sec

3. Storage Estimation (For 5 Years)

- 20 million records per day
- **5 years** of storage
- 1 record = 100 Bytes

Step 1: Total records in 5 years

```
Total Records = 20,000,000 \times 365 \times 5
= 36,500,000,000 (36.5 billion records)
```



```
Total Storage = Total records \times Size per record

= 36.5 \times 10^{9} \times 100

= 3.65 \times 10^{12} Bytes

= 3.65 TB
```

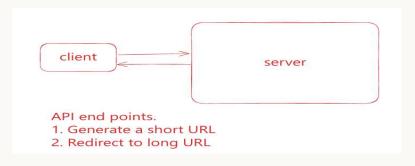
Result:

Storage required ≈ 3.65 TB (for 5 years)



Metric	Value		
Write QPS	≈ 232 writes/sec		
Read QPS	≈ 2,320 reads/sec		
Storage (5 years)	≈ 3.65 TB		

Client ↔ **Server Interaction & API Design**



Task

- Generate a short URL
- Redirect to the long/original URL

APIs Required

- 1. Generate a Short URL (POST Request)
- **Endpoints Options:**
 - POST www.bit.ly/api/v3/generate?longURL=value → (Query parameter)
 - POST www.bit.ly/api/v3/generate/{longURL} → (Path variable)

```
POST www.bit.ly/api/v3/generate
Body: { "longURL": "value" }

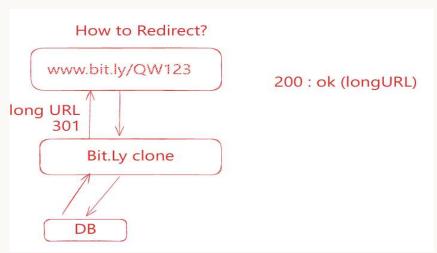
How it works:

Client sends:
{
    "longURL": "https://example.com/very/long/path"
}
    •

Server returns:
{
    "shortURL": "www.bit.ly/Qw23"
```

2. Redirect to Long URL (GET Request)

}



- Endpoint: GET www.bit.ly/{shortURL}
- Server response: Redirect to the original long URL

How Redirection Works

- Client can't redirect using 200 OK that only means success, not redirection
- To redirect, server must use HTTP status codes understood by browsers
- Server hits the database to get the long URL from the short URL

Nedirection Status Codes

301 – Permanent Redirect

- URL: www.bit.ly/Qw23
- Server responds with 301 status code
- Tells browser: this short URL will always go to the same long URL
- Browser does client-side caching (like cookies or local storage)
- On next visit: browser doesn't ask server it directly redirects

♦ 302 – Temporary Redirect

- URL: www.bit.ly/Qw23
- Server responds with **302** status code
- Tells browser: this redirection is temporary
- Browser does not cache the result
- Each time user uses the short URL:
 - It hits the server again
 - o Server can send updated or different long URL if needed

- Use 301 when redirection is fixed & won't change
- Use 302 when redirection might change in future

- 302 is useful when you want to:
 - Track every hit
 - Change long URL later
 - Avoid client-side caching

301 vs 302 Redirect – Kab Kaunsa Use Karein?

301 – Permanent Redirect

✓ Flow:

User \rightarrow bit.ly/Qw23 \rightarrow (Browser handles redirect) \rightarrow Long URL

Pros:

- **V** Browser cache karta hai → No repeated server calls
- Less load on server/database
- Fast redirection for users

X Cons:

- X No tracking (server ko hit hi nahi hota)
- X Cannot update long URL later
- 302 Temporary Redirect
- ✓ Flow:

User \rightarrow bit.ly/Qw23 \rightarrow (Server handles redirect) \rightarrow Long URL

Pros:

- **Track** every request (for analytics)
- Change long URL in future if needed
- Useful for marketing, A/B testing, user behavior tracking

X Cons:

- **X** Browser doesn't cache → **Every visit hits server**
- X Needs scalable backend
- ✓ Recommendation: Use 302 Redirect

Because it gives **flexibility**, **metrics**, and **future control** over URLs.

X System Design: Initial Flow Breakdown

♦ Task 1: URL Shortening

```
\bigcirc Flow:
Long URL \rightarrow [Service] \rightarrow Short URL
```

API:

POST www.bit.ly/api/v3/generate

```
Request Body:
{
    "longURL": "https://example.com/very/long/path"
}
```

```
Response:
{
    "shortURL": "www.bit.ly/Qw23"
}
```

How to Store Mapping Internally

• Use HashMap or HashTable

```
Key → shortURL
Value → longURL

**Example:
task.get("Qw23") → "https://example.com/very/long/path"
```

♦ Task 2: Redirect to Long URL

N Flow:

Short URL \rightarrow [Service] \rightarrow Redirect to Long URL

(IIII) API:

GET www.bit.ly/{shortURL}

- Backend Steps:
 - Extract {shortURL} from request
 - 2. Lookup shortURL in HashMap
 - 3. Fetch corresponding longURL
 - 4. Respond with HTTP 302 Redirect

A Response Example:

HTTP/1.1 302 Found

Location: https://example.com/very/long/path

Important Concepts to Remember

Feature	301 Redirect	302 Redirect
Redirection Type	Permanent	Temporary
Can Track Users?	X No	✓ Yes
Can Update URL?	X No	✓ Yes
Browser Caching	✓ Yes	X No
Use Case	SEO, Static URLs	Analytics, A/B testing, Dynamic URLs

How to Generate a Short URL from a Long URL

Concept: Hash Function-Based URL Shortening

Process:

```
Long URL \rightarrow Hash Function \rightarrow Short URL
```

Example:

Input: https://example.com/very/long/path

Hash: e.g. MD5, SHA-1, SHA-256

Output: bit.ly/a1B2c

HashMap Storage Strategy

Key (Short URL)

P Value (Long URL)

bit.ly/a1B2c

https://example.com/very/long/path

This mapping is stored in the database

X URL Shortening Phase (Storing the URL)

Flow:

```
[User sends Long URL]
[Apply Hash Function \rightarrow Get Short URL]
[Store in HashMap / DB: short \rightarrow long]
[Return Short URL to user]
```

Technologies:

• Use cryptographic hash functions:

MD5, SHA-1, SHA-256 (to ensure uniqueness)



URL Redirection Phase (Reading from DB)

G Flow:

```
[User hits Short URL]
[Lookup: map.get(shortURL)]
[Fetch Long URL]
[Redirect user to Long URL]
```

Problem: In-Memory HashMap Is Not Scalable

X Limitations:

- Stored in RAM
- Works only on a single server
- As data grows (millions/billions of URLs), memory will overflow
- Cannot scale horizontally

Scalable Solution: Use a NoSQL Database

Use:

- Amazon DynamoDB (Key-Value store) recommended
- **MongoDB** (Document store) more complex
 - Replace in-memory HashMap with persistent, scalable DB

High-Level Design (HLD): Distributed System

Architecture Components:

- Multiple Web Servers (for high availability)
- Multiple Database Servers
- **DB Replication** (for fault tolerance)

• Sharding (to distribute data)

Problem in Sharding: Which DB Stores Which URL?

✓ Use: Consistent Hashing

Now We Need Two Hash Functions:

Purpose	Hash Function	Role
To find which DB shard to store in	SHA-1	Distribute data across DBs
To generate the short URL	MD5	Create a unique short key

Full Flow:

```
[Long URL] \downarrow
1. Apply SHA-1 \rightarrow Identify DB shard
2. Apply MD5 \rightarrow Generate short URL \downarrow
[Store shortURL \rightarrow longURL in selected DB]
```

- This ensures:
- Efficient load balancing
- Consistent short URL generation
- Horizontally scalable architecture

Summary

Ste p	Action	Tool
1	Short URL Generation	MD5 / SHA-256
2	Sharding / DB Selection	SHA-1 + Consistent Hashing
3	Storage	NoSQL DB (e.g., DynamoDB)

- Redirect Lookup
 - shortURL → longURL mapping

- 5
- Scaling

Distributed Servers + DB Replication

Better Design – NoSQL → SQL DB

Why Move from NoSQL to SQL?

Client → Server → Short URL

Step 1: Short URL Generation (Initial Logic – NoSQL)

(1) Long URL \rightarrow Hash Function \rightarrow Short URL

DB Used: DynamoDB (NoSQL)

Short URL Long URL

short long

A Problem:

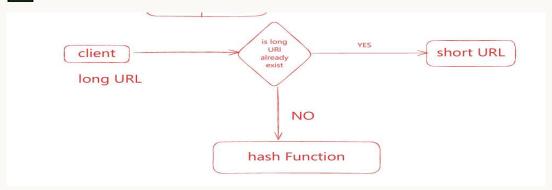
- अगर user same Long URL दोबारा भेजे
- और POST request करे कि "मुझे फिर से short URL दो"
- तो:

Long URL \rightarrow Same hash function \rightarrow Same output

- लेकिन system logic फिर से hash function apply करेगा
- और same **या दूसरा** short URL बना देगा
- जो पहले ही generate किया जा चुका है

🔁 Same hash function हमेशा same number generate करेगा लेकिन अगर check नहीं किया गया कि long URL already है, तो duplicate entry बन सकती है

Solution:



Improved Flow:

```
Client \rightarrow Long URL \rightarrow Check: Already present?
If YES \rightarrow Return existing Short URL
If NO \rightarrow Apply hash function \rightarrow Generate Short URL \rightarrow Store in DB
```

- इससे बार-बार same Long URL पर duplicate short URLs नहीं बनते
- DB में पहले से मौजूद Short URL को ही reuse किया जाता है

Problem in Earlier NoSQL Logic:

• पुरानी mapping:

```
Key → Short URL
Value → Long URL
```

लेकिन अगर आपको check करना है कि Long URL पहले से है या नहीं, तो:

आपको Long URL पर query करनी होगी जो कि NoSQL में मुश्किल है या inefficient हो सकती है



Production-Level Application Requirements:

(1) Query:

For a particular input (user), how many URLs have been generated till now?

SQL supports such **complex queries** efficiently

(2) Performance:

Write operations are faster in SQL compared to NoSQL in this scenario

SQL Table Design (With Indexing)

Column Name Description

ID Primary Key (Auto ID)

Short URL Encoded (shortened)

URL

Long URL Original full URL

Corresponding ID system के लिए indexing useful होगी

$extstyle oldsymbol{ ilde{ ilde{ ilde{P}}}}$ Hash Functions – Long URL o Short URL

Hashing Algorithms:

MD5, SHA-1, SHA-256

Problem:

Hash function se jo value milti hai, wo bahut badi hoti hai, lekin humein short URL chahiye hota hai.

Allowed Characters in Short URL:

 $(0-9) \rightarrow 10$

 $(a-z) \rightarrow 26$

 $(A-Z) \rightarrow 26$

Total = 62 characters

Interview Insight: Keep Short URL "as small as possible"

Example:

www.bit.ly/asd (3 characters)

🔢 Permutations: How Many URLs Can Be Formed?

Characters (n) Possibilities (62ⁿ)

1 char 62

2 chars 3,844

3 chars 238,328

4 chars 14.7 million

5 chars 916 million

6 chars ~56.8 billion

Requirement: Store **36.5 billion URLs** \rightarrow So, take **n = 6**

✓ Format:

www.bit.ly/xxxxxx → 6-character Short URL

Agar n < 6, toh system required number of records store nahi kar paayega.

Hash Function Output Trimming

Problem:

Hash value \rightarrow Long (e.g. ASFWS34233PGWPD3432PQLRS) But we only need **6 characters**

Solution:

Trim first 6 characters Example:

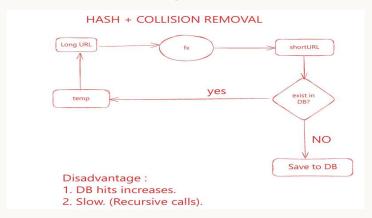
Long Hash \rightarrow ASFWS3



⚠ Collision Possibility:

Same 6-character hash value doosri URL ke liye bhi aa sakti hai.

% Solution: Hashing + Collision Removal



Flow:

```
Long URL
Hash Function \rightarrow 6-char Short URL
Check in DB:
     Yes → Return existing Short URL
     \sqsubseteq No \rightarrow Save to DB
             \downarrow
        If collision:
        - Append temp string (e.g. random or incremental)
        - Re-hash
        - Repeat check
```

Recursive Hashing:

```
Try 1: Hash("www.example.com/page") → ASFWS3 → Exists? → Yes
Try 2: Hash("www.example.com/page1") \rightarrow RTH9KQ \rightarrow Exists? \rightarrow No \rightarrow Save
```

Disadvantages:

More DB Hits: Har collision par DB ko check karna padta hai.

• Slower Performance: Recursive hashing se time badh jaata hai.

BASE-X (Base-62) Approach

Characters Allowed:

Total Characters = 62 Digits: 0-9 (10) Lowercase: a-z (26) Uppercase: A-Z (26)

Base-10 → Base-62 Conversion:

 Decimal
 Base-6

 2

 0-9
 0-9

 10-35
 a-z

 36-61
 A-Z

Example: Convert 14920 into Base-62

Step-by-Step Division:

14920 \div 62 = 240, remainder = 60 \rightarrow Z 240 \div 62 = 3, remainder = 54 \rightarrow S 3 \div 62 = 0, remainder = 3 \rightarrow 3

✓ Final Base-62 Result:

Read from bottom \rightarrow top \rightarrow 3SZ

Mapping in SQL Table:

ID Short URL Long URL

14920 3SZ www.example.com

Advantages of Base-62 Approach

1. No Collisions

(Each ID is unique → Unique Short URL)

X Disadvantages

- 1. Requires a **Unique ID Generator**
- 2. Short URL length is not fixed (as number increases, length increases)
- 3. **Predictable**: Next URL can be guessed (e.g. $3SZ \rightarrow 3T0$)
 - Solution: Use Random Unique ID instead of incremental

Example (Random Unique ID):

Long URL: http://www.example.com/get/code/custID=124

Generated Unique ID: 14920

Short URL = base62(14920) = 3SZ

SQL Table Structure:

ID Short URL Long URL

14920 3SZ www.example.com

Tilde Distributed System Issue

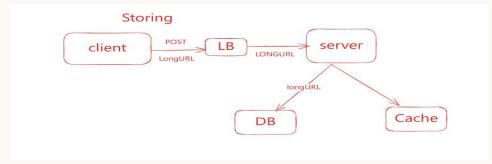
When multiple DBs (shards) are used:

- They may generate the same IDs independently
- Can cause Short URL duplication
 - Solution: Use a Centralized ID Generator

I High-Level System Architecture (HLD)

Storing (POST Flow)

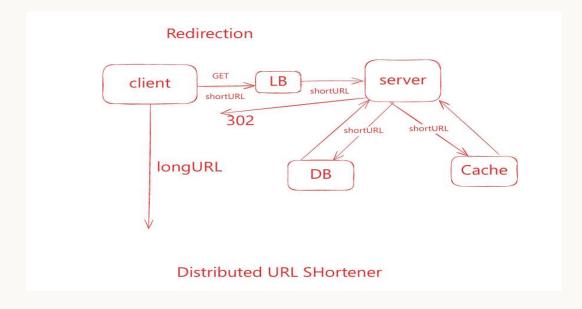
Client
$$\rightarrow$$
 POST \rightarrow Load Balancer \rightarrow Server \downarrow \downarrow Long URL \rightarrow Logging \downarrow DB + Cache



Redirection (GET Flow)

Client
$$\to$$
 GET (Short URL) \to LB \to Server \downarrow Find Long URL \downarrow Response: 302 Redirect \downarrow Client goes to Long URL

302: Temporary redirect → useful for tracking hits



Final Design: Distributed URL Shortener

System supports:

- Centralized ID generation
- Load balancing
- Caching
- Database sharding (with caution)

? Common Interview Question

Q: What if Unique ID generates a value that gives more than 6 characters in Base-62?

A: Not possible.

```
62^6 = 56.8 Billion unique URLs possible \rightarrow We only need to store ~36.5 Billion URLs \rightarrow 6 characters in base-62 are enough
```

Quick Recap:

Requirements

Functional:

- Accept Long URL
- Generate Short URL (e.g., bit.ly/3xYzAb)
- Redirect Short URL to Long URL

Non-Functional:

- Scalable (Millions of URLs/day)
- High Availability
- Low Latency (Fast redirects)

Key Concepts

Metric Value

Write QPS ~232/sec

Read QPS ~2,320/sec

Storage (5 years) ~3.65 TB

Read:Write Ratio 10:1

Characters Allowed 0–9, a–z, A–Z = 62 total

Short URL Length 6 characters (626 ≈

56B)

X URL Generation Methods

♦ 1. Hash Function Approach

- Use: MD5 / SHA-1 / SHA-256 \rightarrow Get hash \rightarrow Trim first 6 chars
- Handle Collisions by:
 - Appending temp strings
 - o Re-hashing until unique
- X Disadvantage: More DB hits, slower due to recursion

♦ 2. Base-62 Encoding (ID → Short URL)

- Convert auto-increment ID to base-62 (e.g., 14920 → 3SZ)
- Collision-free
- X Predictable, variable-length, needs central ID generator

API Design

♦ POST /generate

Request:

```
{ "longURL": "https://example.com/very/long/path" }
Response:
{ "shortURL": "bit.ly/3xYzAb" }

    GET /{shortURL}
```

→ Redirects to original long URL

Redirection: 301 vs 302

Туре	Cache ?	Trackable?	Use When
301	Yes	X No	Permanent redirect (SEO)
302	💢 No	Yes	Analytics, A/B testing
✓ Recommended: 302 → Flexible + Trackable			

Scalable Architecture (HLD)

- 1. Client → Load Balancer → Servers
- 2. Centralized ID Generator
- 3. DB + Caching Layer (e.g., Redis)
- 4. Sharded DB + Consistent Hashing for Long URL storage
- 5. Use NoSQL (DynamoDB) or SQL (MySQL/PostgreSQL) as needed

Optimization Notes

- Use HashMap in-memory only for prototype
- Use NoSQL for fast writes, SQL for analytics/queries
- Handle same long URL with **deduplication logic**

- Use **Base62** for compact, user-friendly URLs
- 6-char short URL supports **up to 56.8 Billion** entries